

AlphaGo

By Philip Dietrich and Laurenz Hemmen

Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion

Structure

1. **The Game of Go**
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion



The Game of Go

- Literally: 'encircling game'
- Two players
- 19x19 fields
- Goal: surround more territory than opponent
- Simple rules, though extremely difficult



The Rules of Go


- Players place stones on vacant points (alternately)
- Adjacent stones of one color form a group
- Liberty of a group: number of vacant adjacent points (group shares liberties)
- If a group has no liberty (gets captured) it is removed from the board



Images from:
https://en.wikibooks.org/wiki/Go/Lesson_2:_Basic_Rules_and_Foundational_concepts

The Rules of Go

- Previous board positions are not allowed to be repeated (Ko-Rule)
- score = number of stones on board + intersections surrounded by stones
- Players can skip
- Game ends if both players skip/ one player gives up



Structure

1. The Game of Go
- 2. DeepMind's AlphaGo**
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion

DeepMind's AlphaGo - Go vs. Chess

	Go	Chess
approx. #(possible moves)	250	35
#(moves in a game)	150	80
Estimated number of different games	10^{761}	10^{120}

Monte Carlo Tree Search

- Used by strong amateur level Go AIs (Fuego, Pachi, Zen, ...)
- Basic Idea:
 - randomly simulate many games (to the end)
 - remember how often node has been visited, how often has that led to a win
 - direct random selection to favor nodes that led previously to wins
- Compared to Deep Blue no knowledge of game necessary

AlphaGo

- SL (rollout) policy network
 - Networks trained on human expert positions to predict next human move
 - One fast softmax network, one big convolutional network (**later more!**)
- RL policy network
 - Playing against different versions of itself
- Value network
 - Predicting probability of a Win
- Tree Search
 - Combining outputs of networks to evaluate value of position
 - Rollout network used to traverse game tree

Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. **Data**
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion

Dataset




- Human Go games from KGS Go Server from games between 2001 and 2007
- Roughly 2.000.000 positions for training (95%) and validation (5%)
- 8000 position for final testing

Representation of data: as feature planes

- Encoding Information on feature planes
 - Positions of own/opponent's stones
 - Liberties of stones
 - if move would be illegal due to ko rule
 - Special patterns
 - Distance to last move, ...
- One plane encodes one of the information for every field on the board

Representation of data: as image

- Own stone: 1, opponent stone 0, vacant field 0.5



Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. **Softmax Network**
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion

Softmax Network: Features and Architecture

Features are all binary for each of the 19x19 squares:

- Colors of stones
- Liberties of stones
- Forbidden fields due to Ko rule
- Neighbors to last move
- Certain 3x3 patterns around candidate moves


⇒ Input size: $19 \times 19 \times (\text{\#feature planes})$

⇒ Output size: $19 \times 19 = 361$

⇒ Number of parameters: $>1.000.000$ for 8 feature planes

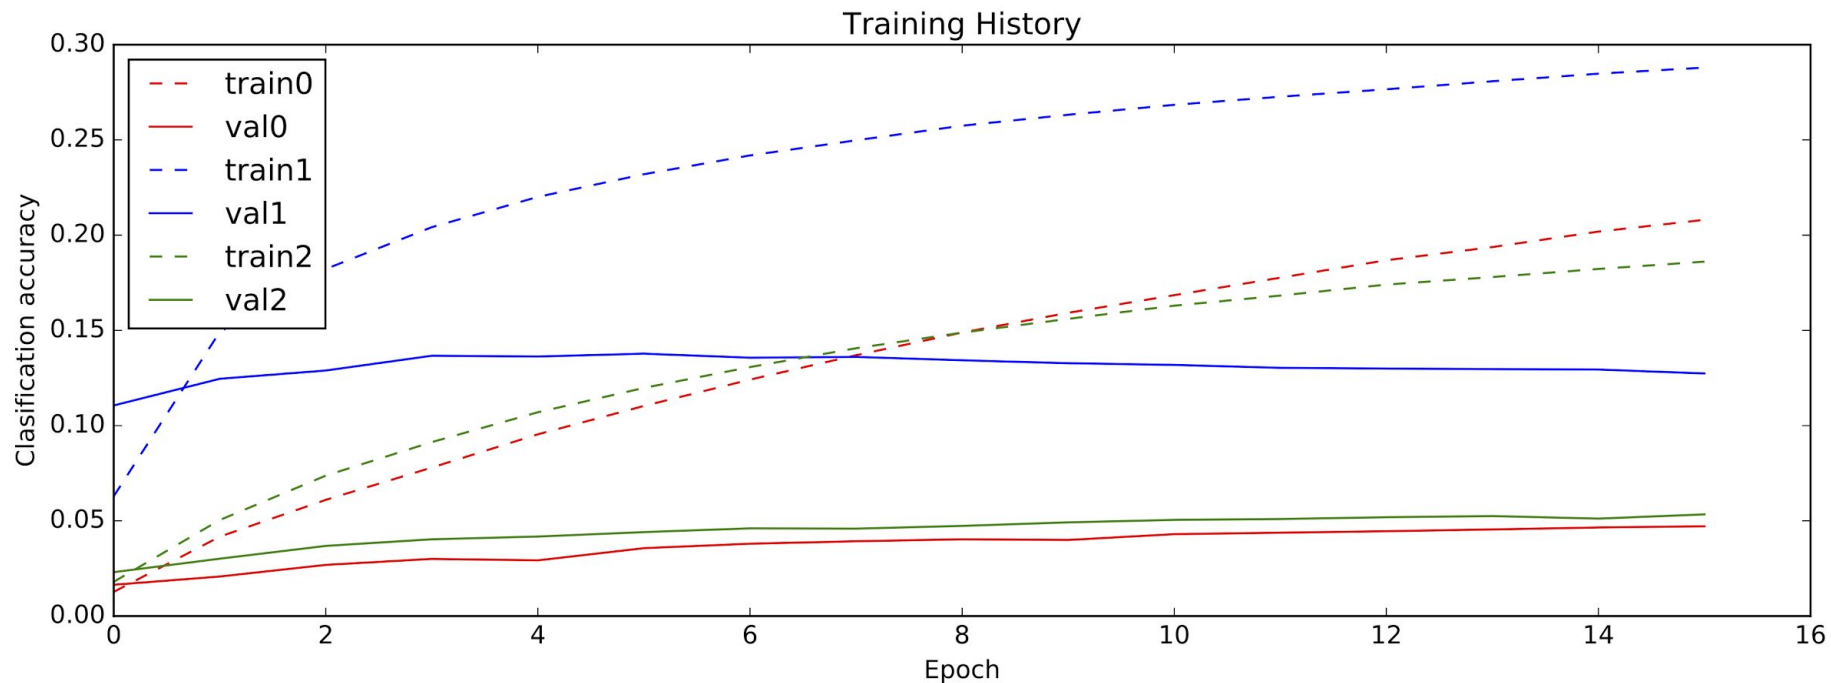
Softmax: Activation function

Reminder:




For small and huge values, gradient nearly vanishes during backprop \Rightarrow no/small parameter update

Results: Softmax networks



Results: Softmax networks




Softmax: Different features compared

	Test Accuracy [%]	time/prediction [ms]
Col+Lib, ko	6.7	0.11
Col, Lib	7.1	0.10
Col, Lib, Neighbors	17.1	0.13
Neighbors	6.7	0.11
AlphaGo's Softmax	24.2	-
Baseline: Guessing	>0.3 (=100/361)	-

Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
- 5. Convolutional Neural Network**
 - a. Architecture
 - b. Results
6. Comparing Models
7. Conclusion

Convolutional network




Taken from:

<https://developer.apple.com/library/content/documentation/Performance/Conceptual/vImage/ConvolutionOperations/ConvolutionOperations.html>

Convolutional network



- Moves filter over image pixels
- Certain shapes give stronger output
- Normally several filters applied to one image
- Hyperparameters: number and size of filters
- Parameter determined by backpropagation: weights of filters



Taken from:
http://blog.csdn.net/stdcoutz_yx/article/details/42091205

Convolutional network


- So far: fully convolutional, in AlphaGo: additional nonlinearities
- Biological view: (taken from <http://cs231n.github.io/neural-networks-1/#bio>)



Convolutional network

Activation function used here: Rectified Linear Unit (ReLU)

- Faster to compute than tanh, sigmoid
- No saturation in positive spectrum
- Threshold at zero



taken from:

<http://cs231n.github.io/neural-networks-1/>
#bio

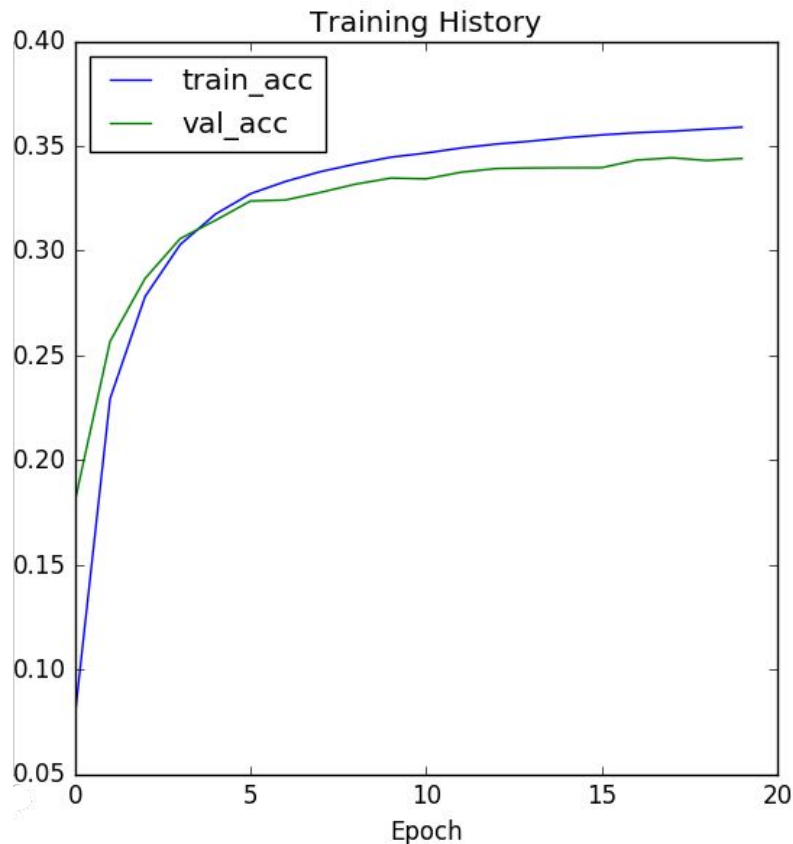
Convolutional Networks

Network		Accuracy one board layer [%]	Accuracy 3 board layers [%]	Accuracy 7 board Layers [%]	Accuracy Col + Lib, Ko [%]
Convolutional - ReLu - Softmax	4 filters	4.0	10.3	-	6.7
	32 filters	3.8	14.2	26.4	7.2
Convolutional - Softmax	4 filters	1.3	-	-	7.8
	32 filters	1.5	-	-	8.5
AlphaGo	192 filters/layer	-	55.7	-	57.0

Convolutional Networks: More Conv Layers

Approach from DeepMind Paper:

- Convolutional Layer with k 5x5 filters
 - Convolutional Layer with k 3x3 filters
 - Convolutional Layer with 1 1x1 filter
 - Softmax-Layer
-
- AlphaGo: k = 196, we: k = 32
 - Number of parameters ≈ 145.000
 - \Rightarrow no overfitting!
 - 7 input layers: 34.4 % accuracy
 - 3 input layers: 17.9 % accuracy



Convolutional Networks: More Conv Layers

- Even more layers:
 - 6 Convolutional layers (48 * 5x5, 32*5x5, 32* 5x5, 32*3x3, 8*3x3, 1*1x1)
 - ReLu activation function on every layer
 - One softmax layer in the end
- 214763 Parameters
- Used 1.000.000 training and validation samples
- Final test accuracy: 39.1% (training accuracy 43.0%)

⇒ best accuracy we could get

Convolutional networks

- Many more variations of #filters, sizes of filters,... tested
 - More filters [4, 8, 16, 32] result in higher accuracy, but more overfitting
 - Filter sizes [3x3, 5x5, 7x7] cause nearly no difference, 5x5 slightly better
 - More training data resulted in significantly higher accuracies
- Only small variations, results at similar accuracies
- General impression: the more layers used, ...
 - the longer it took to train the network
 - the easier it was to overfit the data to to high number of parameters \Rightarrow more data needed than we actually had (AlphaGo: 30 Mio, we: 2 Mio)
 - the longer the prediction time

Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. **Comparing Models**
7. Conclusion

Comparison: Softmax vs. ConvNet

	Convolutional Networks							Softmax		
	Conv - ReLU - SM			Conv - SM		CRCR CRS	(CR) ⁶ S			
	8p, 4f	8p, 32f	3p, 32f	8p, 4f	8p, 32f	7p, 32f	13p,48f	6 p	7 p	8 p
Runtime [ms]	0.24	0.24	0.24	0.13	0.21	0.35	0.54	0.06	0.06	0.08
Accuracy [%]	6.7	7.2	14.2	7.8	8.5	34.4	39.1	7.1	17.1	6.7

Structure

1. The Game of Go
2. DeepMind's AlphaGo
3. Data
 - a. Representations of Boards
4. Softmax Network
 - a. Architecture
 - b. Results
5. Convolutional Neural Network
 - a. Architecture
 - b. Results
6. Comparing Models
7. **Conclusion**

Conclusion: Problems during Training

- Overfitting to training data due to high number of parameters in model compared to number of training positions
- memory/ processing limitations
- Very long training times (often: training time > 10 hours)

Conclusion: How well does it play?

- Doesn't "know" the rules, can't pass \Rightarrow can't really play
- Only tries to predict next human move, not necessarily best move \Rightarrow AlphaGo uses additional RF Learning model playing against versions of itself
- Even if most moves are predicted correctly, false predictions can give opponent opportunity to easily win the game
- May create uncommon positions, isn't trained on these

Conclusion: Possible further improvements

- Symmetry
 - Using mirrored board positions
 - Forcing weights to be symmetrical
- Add Go rules and already occupied fields
- Add additional layers (Dropout, more Convolutions, ..)
- Use more data (we: 2e6 vs. AlphaGo: 30e6 positions)
- Additional information / feature planes
 - Possible captures, stone saves
 - Matching certain kind of patterns
 - around last move
 - Around every position

Conclusion: Possible further improvements

Extended Data Table 4 | Input features for rollout and tree policy

Feature	# of patterns	Description
Response	1	Whether move matches one or more response pattern features
Save atari	1	Move saves stone(s) from capture
Neighbour	8	Move is 8-connected to previous move
Nakade	8192	Move matches a <i>nakade</i> pattern at captured stone
Response pattern	32207	Move matches 12-point diamond pattern near previous move
Non-response pattern	69338	Move matches 3×3 pattern around move
Self-atari	1	Move allows stones to be captured
Last move distance	34	Manhattan distance to previous two moves
Non-response pattern	32207	Move matches 12-point diamond pattern centred around move

Features used by the rollout policy (first set) and tree policy (first and second set). Patterns are based on stone colour (black/white/empty) and liberties (1, 2, ≥ 3) at each intersection of the pattern.

Conclusion

- Our network wasn't nearly as good as AlphaGo's
- Convolutional networks gave better results
- Softmax networks were faster
- Best network architecture: some Convolutional layers (decreasing in size, #filters) , finally one softmax layer
 - Relatively low number of parameters (in the order $1e6$) \Rightarrow data less overfitted
 - Better results than using fully connected layers

Questions?

Bibliography

- [1] Silver, Huang, Maddison, Guez; Mastering the game of Go with deep neural networks and tree search
- [2] Christopher Clark, Amos Storkey: Training Deep Convolutional Neural Networks to Play Go